

Day 3: Parallel Computing

Thomas Hauser
Director of Research Computing
University of Colorado Boulder
thomas.hauser@colorado.edu

Outline

- Topics today
 - Computer Architecture 101
 - Single processor optimization
 - Shared memory parallel programming
 - Programming accelerators
 - Distributed memory parallel programming
- Provide you with
 - Overview and simple guidelines for code development
 - Simple examples
 - Pointers to other materials

References

1. Hager, G, and G Wellein. 2010. *Introduction to High Performance Computing for Scientists and Engineers*. CRC Press, Jul 2, 2010 - 356 pages.
2. Levesque, John, and Gene Wagenbreth. 2010. *High Performance Computing*. CRC Press.
3. Stampede User Guide:
<https://www.tacc.utexas.edu/user-services/user-guides/stampede-user-guide>
4. Henry Neeman, Supercomputing in Plain English, A High Performance Computing Workshop Series,
<http://www.oscer.ou.edu/education.php>

Computer Architecture 101

Stampede Supercomputer




Stampede has 6,400 nodes [3]
56 GB/s FDR Infiniband interconnect



www.scan.co.uk

Socket:

- 2.7 GhZ
- 8 Cores
- 8 DP FP operations per clock cycle
- 64 GB L1 Cache/core
- Vector width: 4 double precision items




Stampede compute node [3]:

- 2 Sockets per Node → 2 Xeon E5 processors
- 1 Xeon Phi coprocessor
- 32 GB Memory
- 250 GB Disk

Parallelism at All Levels


- Parallelism across multiple nodes or process
- Parallelism across threads
- Parallelism across instructions
- Parallelism on data
 - SIMD Single Instruction Multiple Data



www.scan.co.uk

Von Neumann architecture

- Instruction and data are stored in memory
- Instructions are read by a control unit
- Arithmetic/logic unit is responsible for computation
- CPU consists of control and arithmetic units with interfaces to memory and I/O
- Compiler translates high level language into instructions that can be stored and executed



Difficulties with the van Neumann architecture

- Instructions and data must be fed continuously to control and arithmetic units
 - Speed of memory is the limiting factor on compute performance
 - *Neumann bottleneck*
- Sequential architecture
 - SISD (Single Instruction Single Data)
- Modified and extended for modern processors
 - Parallelism on all levels of a machine

Performance Metrics

- Every component of a CPU can operate at some maximum speed
 - **Peak Performance**
- Applications generally cannot achieve peak performance
- Floating-point operations per second (FLOPS)
 - Counting only multiply add
 - Other operations are more expensive
- Modern processors can deliver at between 2 or 8 FLOPS per clock cycle
- Typical clock frequency between 2 and 3 GHz

Floating Point Performance

$$P = n_{\text{core}} * F * S * \nu$$

- Example: Intel Xeon E5 on Stamped
 - Number of cores: 8 n_{core}
 - FP instructions per cycle: 2 (1 Multiply and 1 add) F
 - FP operations / instruction (SIMD): 4 (dp) / 8 (sp) S
 - Clock speed: 2.7 GHZ ν

$$P = 173 \text{ GF/s (dp)} \quad \text{or} \quad 346 \text{ GF/s (sp)}$$

- But: P= 5.4 GF/s (dp) for serial, non-SIMD code

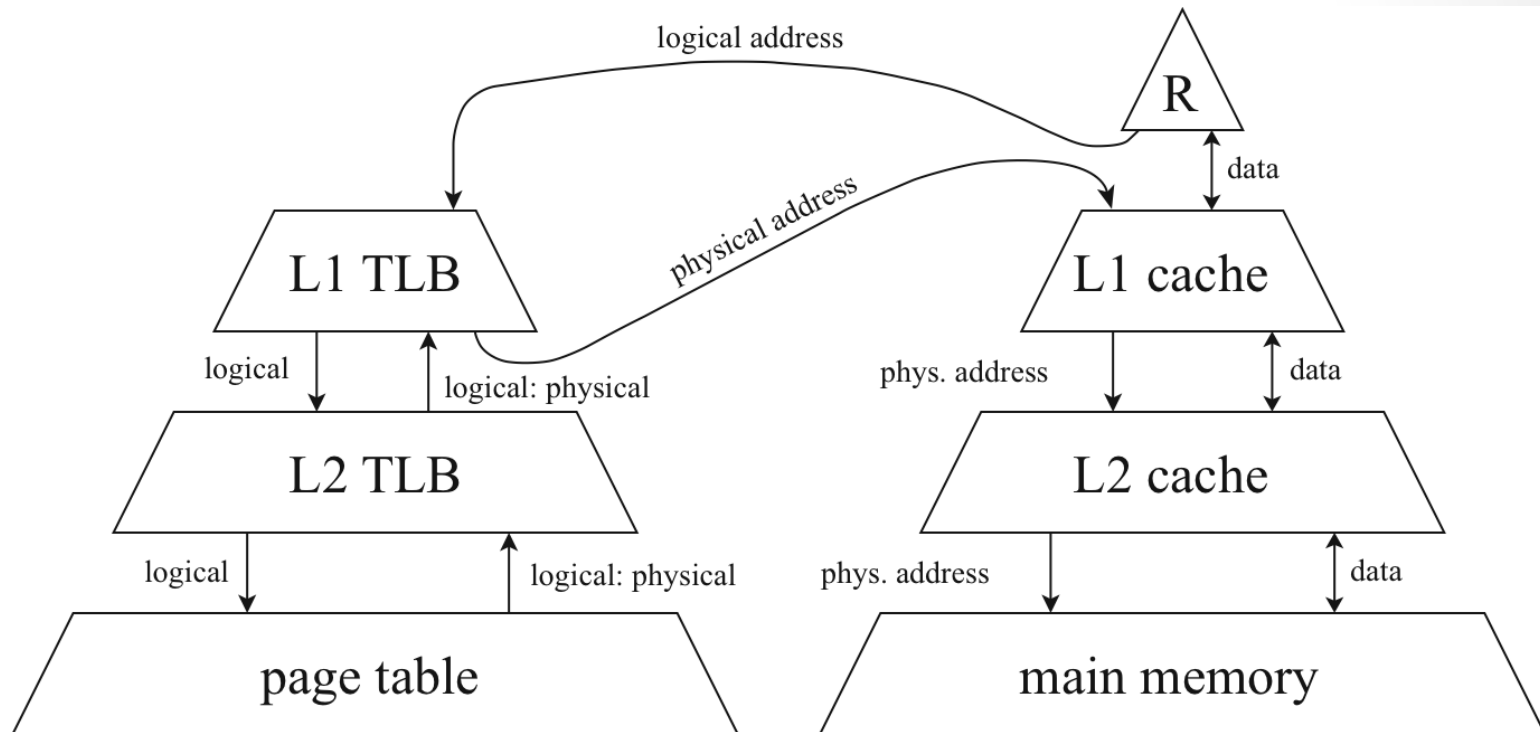
Memory Performance

- Path from and to caches and memory is generally the bottle neck for application performance.
- Bandwidth is the performance measure for memory
 - Gbytes/sec

Data-centric view of a microprocessor

- GFlops/sec
- Gbytes/sec
- Describe the most relevant performance features of microprocessors
- Memory access is generally the limiting factor for performance
 - Data centric view

Memory system

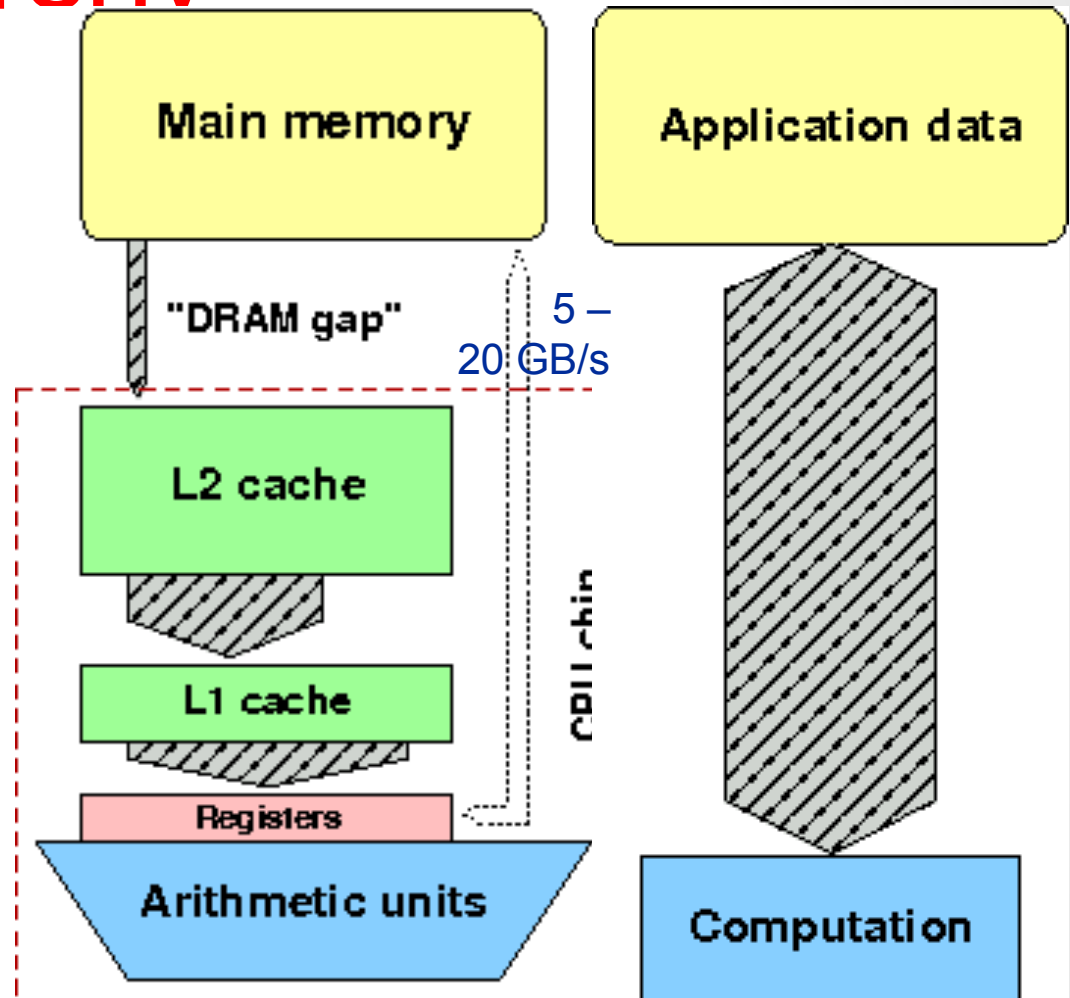


- R: register
- Right: Data and instruction caches
- Left: Translation caches

Memory hierarchy

1. CPU/Arithmetic unit issues a load request to transfer a data item to a register
2. Cache logic automatically checks all cache levels if data item is already in cache
3. If data item is in cache “**cache hit**” it is loaded to register.
4. If data item is in no cache level (“**cache miss**”) data item is loaded from main memory and a copy is held in cache

If cache is already full another cache line must be invalidated or “evicted” in 4



Optimize for memory access

- Page Fault, file on IDE disk: 1.000.000.000 cycles
- Page Fault, file in buffer cache: 10.000 cycles
- Page Fault, file on ram disk: 5.000 cycles
- Page Fault, zero page: 3.000 cycles
- Main memory access: about 200 cycles
- L3 cache hit: about 52 cycles
- L1 cache hit: 2 cycles

The Core i7 can issue 4 instructions per cycle. So a penalty of 2 cycles for L1 memory access means a missed opportunity for 7 instructions.

Peak vs. Realized Performance

Peak performance = guaranteed not to exceed

Realized performance = what you achieve

Scientific applications realize as low as
1-25% of peak on microprocessor-based systems


Reason: mismatch between application and architecture capabilities

- Architecture has insufficient bandwidth to main memory:
 - microprocessors often provide < 1 byte from memory per FLOP
 - scientific applications often need more
- Application has insufficient locality
 - irregular accesses can squander memory bandwidth
 - use only part of each data block fetched from memory
 - may not adequately reuse costly virtual memory address translations
- Exposed memory latency
 - architecture: inadequate memory parallelism to hide latency
 - application: not structured to exploit memory parallelism
- Instruction mix doesn't match available functional units

Low-level benchmarking

- Benchmark to test some specific feature of the architecture
 - Peak performance
 - Memory bandwidth
- Example of benchmark is *vector triad*
 - Measures performance of data transfers between memory and arithmetic units
 - $A(:) = B(:) + C(:) * D(:)$
 - 3 load streams
 - 1 store stream
 - 2 Flops

Results for Triad from my Laptop



Stream example

- See notebook

Performance Analysis and Tuning

- Increasingly necessary
 - Gap between realized and peak performance is growing
- Increasingly hard
 - Complex architectures are harder to program effectively
 - complex processors: pipelining, out-of-order execution, VLIW
 - complex memory hierarchy: multi-level non-blocking caches, TLB
 - Optimizing compilers are critical to achieving high performance
 - small program changes may have a large impact
 - Modern scientific applications pose challenges for tools
 - multi-lingual programs
 - many source files
 - complex build process
 - external libraries in binary-only form
 - Many tools don't help you identify or solve your problem

Tuning Applications

- Performance is an interaction between
 - Numerical model
 - Algorithms
 - Problem formulation (as a program)
 - Data structures
 - System software
 - Hardware
- Removing performance bottlenecks may require dependent adjustments to all

Having fun with compiler options

The Joy of Compiler Options

- Every compiler has a different set of options that you can set.
- Among these are options that control single processor optimization: superscalar, pipelining, vectorization, scalar optimizations, loop optimizations, inlining and so on.

GCC optimization options

-falign-functions[=*n*] -falign-jumps[=*n*] -falign-labels[=*n*] -falign-loops[=*n*] -fassociative-math -fauto-inc-dec -fbranch-probabilities -fbranch-target-load-optimize -fbranch-target-load-optimize2 -fbtr-bb-exclusive -fcaller-saves -fcheck-data-deps -fcombine-stack-adjustments -fconserve-stack -fcompare-elim -fcprop-registers -fcrossjumping -fcse-follow-jumps -fcse-skip-blocks -fcx-fortran-rules -fcx-limited-range -fdata-sections -fdce -fdce -fdelayed-branch -fdelete-null-pointer-checks -fdse -fdvirtualize -fdse -fearly-inlining -fipa-sra -fexpensive-optimizations -ffast-math -ffinite-math-only -ffloat-store -fexcess-precision=*style* -fforward-propagate -ffp-contract=*style* -ffunction-sections -fgcse -fgcse-after-reload -fgcse-las -fgcse-lm -fgraphite-identity -fgcse-sm -fif-conversion -fif-conversion2 -findirect-inlining -finline-functions -finline-functions-called-once -finline-limit=*n* -finline-small-functions -fipa-cp -fipa-cp-clone -fipa-matrix-reorg -fipa-pta -fipa-profile -fipa-pure-const -fipa-reference -fipa-struct-reorg -fira-algorithm=*algorithm* -fira-region=*region* -fira-loop-pressure -fno-ira-share-save-slots -fno-ira-share-spill-slots -fira-verbose=*n* -fivopts -fkeep-inline-functions -fkeep-static-consts -floop-block -floop-flatten -floop-interchange -floop-strip-mine -floop-parallelize-all -flto -flto-compression-level -flto-partition=*alg* -flto-report -fmerge-all-constants -fmerge-constants -fmodulo-sched -fmodulo-sched-allow-regmoves -fmove-loop-invariants -fmudflap -fmudflapir -fmudflapth -fno-branch-count-reg -fno-default-inline -fno-defer-pop -fno-function-cse -fno-guess-branch-probability -fno-inline -fno-math-errno -fno-peephole -fno-peephole2 -fno-sched-interblock -fno-sched-spec -fno-signed-zeros -fno-toplevel-reorder -fno-trapping-math -fno-zero-initialized-in-bss -fomit-frame-pointer -foptimize-register-move -foptimize-sibling-calls -fpartial-inlining -fpeel-loops -fpredictive-commoning -fprefetch-loop-arrays -fprofile-correction -fprofile-dir=*path* -fprofile-generate -fprofile-generate=*path* -fprofile-use -fprofile-use=*path* -fprofile-values -freciprocal-math -fregmove -frename-registers -freorder-blocks -freorder-blocks-and-partition -freorder-functions -frerun-cse-after-loop -freschedule-modulo-scheduled-loops -frounding-math -fsched2-use-superblocks -fsched-pressure -fsched-spec-load -fsched-spec-load-dangerous -fsched-stalled-insns-dep[=*n*] -fsched-stalled-insns[=*n*] -fsched-group-heuristic -fsched-critical-path-heuristic -fsched-spec-insn-heuristic -fsched-rank-heuristic -fsched-last-insn-heuristic -fsched-dep-count-heuristic -fschedule-insns -fschedule-insns2 -fsection-anchors -fselective-scheduling -fselective-scheduling2 -fsel-sched-pipelining -fsel-sched-pipelining-outer-loops -fsignaling-nans -fsingle-precision-constant -fsplit-ivs-in-unroller -fsplit-wide-types -fstack-protector -fstack-protector-all -fstrict-aliasing -fstrict-overflow -fthread-jumps -ftracer -ftree-bit-ccp -ftree-builtin-call-dce -ftree-ccp -ftree-ch -ftree-copy-prop -ftree-copyrename -ftree-dce -ftree-dominator-opts -ftree-dse -ftree-forwprop -ftree-fre -ftree-loop-if-convert -ftree-loop-if-convert-stores -ftree-loop-im -ftree-phi-prop -ftree-loop-distribution -ftree-loop-distribute-patterns -ftree-loop-ivcanon -ftree-loop-linear -ftree-loop-optimize -ftree-parallelize-loops=*n* -ftree-pre -ftree-pta -ftree-reassoc -ftree-sink -ftree-sra -ftree-switch-conversion -ftree-ter -ftree-vect-loop-version -ftree-vectorize -ftree-vrp -funit-at-a-time -funroll-all-loops -funroll-loops -funsafe-loop-optimizations -funsafe-math-optimizations -funswitch-loops -fvariable-expansion-in-unroller -fvect-cost-model -fvpt -fweb -fwhole-program -fwpa -fuse-linker-plugin --param *name=value* -O -O0 -O1 -O2 -O3 -Os -Ofast

Example Compile Lines

- GCC
 - `gfortran -Ofast`
- Intel
 - `ifort -Ofast`
- Portland Group f90
 - `pgf90 -fastsse -Mipa=fast`

Scalar Optimizations

- Copy Propagation
- Constant Folding
- Dead Code Removal
- Strength Reduction
- Common Subexpression Elimination
- Variable Renaming
- **Vectorization of loops**
- Not every compiler does all of these, so it sometimes can be worth doing these by hand.

What is Vector or SIMD Code

- Processor can execute one instruction on a few elements
 - SSE: 4 elements at a time
 - AVX: 8 elements at a time
 - MIC: 16 elements at a time
- Stream:
 - Scalar code computes result one element at a time



```
for (j=0; j<STREAM_ARRAY_SIZE; j++)  
    a[j] = b[j]+scalar*c[j];
```

Programming Guidelines

- To vectorize a loop you may need to make simple changes
- Use
 - Simple for loops
 - No if statement in loops
 - No function calls
 - Fortran (or use restrict in C on pointers)

Steady State Heat Distribution Problem

Ice bath



Steam


Solving the Problem

- Underlying PDE is the Poisson equation

$$u_{xx} + u_{yy} = f(x, y)$$

- This is an example of an elliptical PDE
- Will create a 2-D grid
- Each grid point represents value of state state solution at particular (x, y) location in plate

Heart of Sequential Program



Heat equation example

- gfortran 4.6.2
- 2.4. GHz Intel Core 2 Duo
- 1000x2000 problem size with 100 iterations
- -O
 - 2.443s
- -O2
 - 2.452s
- -O3
 - 1.163s
- -Ofast
 - 1.147s

Heat equation example

- -O
 - 2.443s
- -O2
 - 2.452s
- -O3
 - 1.163s
- -Ofast
 - 1.147s
- -Ofast -funroll-all-loops
 - 1.123

Loop Optimizations

- Hoisting Loop Invariant Code
- Unswitching
- Iteration Peeling
- Index Set Splitting
- **Loop Interchange**
- Unrolling
- Loop Fusion
- Loop Fission
- Not every compiler does all of these, so it sometimes can be worth doing some of these by hand.

Loop Carried Dependency

```
DO i = 2, length  
  a(i) = a(i-1) + b(i)  
END DO
```

Here, each iteration of the loop depends on the previous iteration. That is, the iteration **i=3** depends on iteration **i=2**, iteration **i=4** depends on **i=3**, iteration **i=5** depends on **i=4**, etc.

This is sometimes called a **loop carried dependency**.

There is no way to execute iteration **i** until after iteration **i-1** has completed, so this loop can't be parallelized.

Reductions Aren't Dependencies

```
array_sum = 0  
DO i = 1, length  
    array_sum = array_sum + array(i)  
END DO
```

- Other kinds of reductions:
 - product, .AND., .OR., minimum, maximum, index of minimum, index of maximum, number of occurrences of a particular value, etc.
- Reductions are so common that hardware and compilers are optimized to handle them.
- Also, they aren't really dependencies, because the order in which the individual operations are performed doesn't matter.

Why Do We Care?

Loops are the **favorite control structures** of High Performance Computing, because compilers know how to optimize their performance using instruction-level parallelism: superscalar, pipelining and vectorization can give excellent speedup.

Loop carried dependencies affect whether a loop can be parallelized, and how much.

Loop Dependency Example

```
do index = 2, length  
  dst(index) = dst(index-1) + dst(index)  
end do
```

```
do index = 2, length  
  dst(index) = dst(index-1) + src2(index)  
end do
```

```
do index = 2, length  
  dst(index) = src1(index-1) + dst(index)  
end do
```

```
do index = 2, length  
  dst(index) = src1(index-1) + src1(index)  
end do
```

```
do index = 2, length  
  dst(index) = src1(index-1) + src2(index)  
end do
```

Hoisting Loop Invariant Code

Code that doesn't change inside the loop is called loop invariant. It doesn't need to be calculated over and over

Before

```
DO i = 1, n
  a(i) = b(i) + c * d
  e = g(n)
END DO
```

After

```
temp = c * d
DO i = 1, n
  a(i) = b(i) + temp
END DO
e = g(n)
```

Unswitching

```
DO i = 1, n
  DO j = 2, n
    IF (t(i) > 0) THEN
      a(i,j) = a(i,j) * t(i) + b(j)
    ELSE
      a(i,j) = 0.0
    END IF
  END DO
END DO
```

```
DO i = 1, n
  IF (t(i) > 0) THEN
    DO j = 2, n
      a(i,j) = a(i,j) * t(i) + b(j)
    END DO
  ELSE
    DO j = 2, n
      a(i,j) = 0.0
    END DO
  END IF
END DO
```

Before

The condition is
j-independent.

After

So, it can migrate
outside the j loop.

Loop Interchange

Before

```
DO i = 1, ni
  DO j = 1, nj
    a(i,j) = b(i,j)
  END DO
END DO
```

After

```
DO j = 1, nj
  DO i = 1, ni
    a(i,j) = b(i,j)
  END DO
END DO
```

Array elements **a(i,j)** and **a(i+1,j)** are near each other in memory, while **a(i,j+1)** may be far, so it makes sense to make the **i** loop be the inner loop. (This is reversed in C, C++ and Java.)

Summary

- Find optimized libraries first
- Try the compiler next
 - Select options
 - Review compiler output, e.g a loop can't be vectorized
- Profile
 - Rewrite only code which is relevant to your performance
- Rewrite code
 - Get the loop ordering right
 - Avoid if statement in a loop
 - Look at cache behavior too
 - Blocking/Tiling

Summary

- Readable and correct code is more important than fast and obfuscated code
- Optimize for memory hierarchy
- Profile to identify hotspots
- Applied to real CFD code

