

Basics of parallelism

Basic approaches to parallelism
Speedup and Efficiency

Material in this presentation from textbook:

Georg Hager and Gerhard Wellein, Introduction to High Performance Computing for Scientists and Engineers, Chapman & Hall/CRC Computational Science Series ISBN 978-1-4398-1192-4

Learning Objectives

- Basic approaches to parallelize problems
- Predict performance of parallel programs
- Understand barriers to higher performance

Outline

- Why parallelize?
- Parallelism
- Parallel scalability
 - Limitations
 - Scalability metrics
 - General speedup formula
 - Amdahl's Law
 - Gustafson-Barsis' Law
 - Karp-Flatt metric

Why parallelize

- Single core too slow for solving the problem in a “reasonable” time
 - “Reasonable” time: overnight, over lunch, duration of a PhD theses
- Memory requirements
 - Larger problem
 - More physics
 - More particles

Parallelism

- For multi-core or multi-node computers
- Data parallelism
 - Single Program Multiple Data (SPMD)
 - Same code is executed on all processors
 - Data is different on the nodes
- Functional parallelism
 - Splitting problem in separate subtasks
 - Multiple Program Multiple Data (MPMD)
 - Subtask could be SPMD
 - Difficult to load balance if subtasks have different performance properties

Examples Data Parallelism

P1	<pre>do i=1,500 a(i)=c*b(i) enddo</pre>	<pre>do i=1,1000 a(i)=c*b(i) enddo</pre>
P2	<pre>do i=501,1000 a(i)=c*b(i) enddo</pre>	

Speedup Formula

$$\text{Speedup} = \frac{\text{Sequential execution time}}{\text{Parallel execution time}}$$

Execution Time Components

- Inherently sequential computations: $s(n)$
- Potentially parallel computations: $p(n)$
- Communication operations: $c(n, p)$
- Speedup expression:

$$S \leq \frac{s+p}{s(n)+p/N+c}$$

Parallel Overhead

- Overhead because of
 - Startup time
 - Synchronizations
 - Communication
 - Overhead by libraries, compilers
 - Termination time

Efficiency

$$\text{Efficiency} = \frac{\text{Sequential execution time}}{\text{Processors} \times \text{Parallel execution time}}$$

$$\text{Efficiency} = \frac{\text{Speedup}}{\text{Processors}}$$

Amdahl's Law

- How much can we improve the time to solution
 - Fixed problem size

$$T_f^s = s + p$$

$$T_f^p = s + \frac{p}{N}$$


$$S = \frac{1}{s + \frac{1-s}{N}}$$

- How much can the serial portion be if we want a speedup of 90 on 100 processors

$$90 = \frac{1}{s + \frac{1-s}{100}} \rightarrow s = 0.0001$$

Limitations of Amdahl's Law

- Ignores overhead
- Overestimates speedup achievable




<http://www.drdobbs.com/parallel/amdahls-law-vs-gustafson-barsis-law/240162980?pgno=2>

Amdahl Effect

- Typically overhead has lower complexity than parallel work
- As *problem size* increases, parallel work dominates overhead
- As n increases, speedup increases

Illustration of Amdahl Effect




Review of Amdahl's Law

- Treats problem size as a constant
- Shows how execution time decreases as number of processors increases
- Strong scaling
 - Problem size is fixed
 - Number of processor increases

Another Perspective

- We often use faster computers to solve larger problem instances
- Let's treat time as a constant and allow problem size to increase with number of processors
- "...speedup should be measured by scaling the problem to the number of processors, not by fixing the problem size" – John Gustafson

Gustafson-Barsis's Law



$$SS(N) = \frac{s+p*N}{s+p} = N + (1 - N)s$$

<http://www.drdoobs.com/parallel/amdahls-law-vs-gustafson-barsis-law/240162980?pgno=2>

Gustafson-Barsis's Law

- Begin with parallel execution time
- Estimate sequential execution time to solve same problem
- Problem size is an increasing function of N
- Predicts **scaled speedup**
- Weak Scaling

Example 1

- An application running on 10 processors spends 3% of its time in serial code. What is the scaled speedup of the application?

$$S = 10 + (1 - 10)(0.03) = 10 - 0.27 = 9.73$$



...except 9 do not have to execute serial code

Execution on 1 CPU takes 10 times as long...

Example 2

- What is the maximum fraction of a program's parallel execution time that can be spent in serial code if it is to achieve a scaled speedup of 7 on 8 processors?

$$90 = 100 + (1 - 100)s \Rightarrow s \approx 0.1$$

The Karp-Flatt Metric

- Amdahl's Law and Gustafson-Barsis' Law ignore overhead
- They can overestimate speedup or scaled speedup
- Karp and Flatt proposed another metric

Experimentally Determined Serial Fraction

$$e = \frac{s+c}{s+p}$$

Inherently serial component
of parallel computation +
processor communication and
synchronization overhead

Single processor execution time

$$e = \frac{1/S - 1/p}{1 - 1/p}$$

Experimentally Determined Serial Fraction

- Takes into account parallel overhead
- Detects other sources of overhead or inefficiency ignored in speedup model
 - Process startup time
 - Process synchronization time
 - Imbalanced workload
 - Architectural overhead

Example 1

p	2	3	4	5	6	7	8
ψ	1.8	2.5	3.1	3.6	4.0	4.4	4.7

What is the primary reason for speedup of only 4.7 on 8 CPUs?

e	0.1	0.1	0.1	0.1	0.1	0.1	0.1
---	-----	-----	-----	-----	-----	-----	-----

Since e is constant, large serial fraction is the primary reason.

Example 2

p	2	3	4	5	6	7	8
ψ	1.9	2.6	3.2	3.7	4.1	4.5	4.7

What is the primary reason for speedup of only 4.7 on 8 CPUs?

e	0.070	0.075	0.080	0.085	0.090	0.095	0.100
-----	-------	-------	-------	-------	-------	-------	-------

Since e is steadily increasing, overhead is the primary reason.

Pop Quiz

p	4	8	12
ψ	3.9	6.5	?

- Is this program likely to achieve a speedup of 10 on 12 processors?

Summary (1/3)

- Performance terms
 - Speedup
 - Efficiency
- Model of speedup
 - Serial component
 - Parallel component
 - Communication component

Summary (2/3)

- What prevents linear speedup?
 - Serial operations
 - Communication operations
 - Process start-up
 - Imbalanced workloads
 - Architectural limitations

Summary (3/3)

- Analyzing parallel performance
 - Amdahl's Law
 - Gustafson-Barsis' Law
 - Karp-Flatt metric