

OpenMP programming


Thomas Hauser
Director Research Computing
thomas.hauser@colorado.edu

Outline

- OpenMP
- Shared-memory model
- Parallel `for` loops
- Declaring private variables
- Critical sections
- Reductions

Parallelism at All Levels

- Parallelism across multiple nodes or process - **MPI**
- Parallelism across threads - **OpenMP**
- Parallelism across instructions
- Parallelism on data – SIMD Single Instruction Multiple Data




www.scan.co.uk

OpenMP

- OpenMP: An application programming interface (API) for parallel programming on multiprocessors
 - Compiler directives
 - Library of support functions
 - Environment variables
- OpenMP works in conjunction with Fortran, C, or C++
- <http://openmp.org>

Shared-memory Model




Processors interact and synchronize with each other through shared variables.

Fork/Join Parallelism

- Initially only master thread is active
- Master thread executes sequential code
- Fork: Master thread creates or awakens additional threads to execute parallel code
- Join: At end of parallel code created threads die or are suspended

Fork/Join Parallelism



Shared-memory Model vs. Message-passing Model (#1)

- Shared-memory model
 - Number active threads 1 at start and finish of program, changes dynamically during execution
- Message-passing model
 - All processes active throughout execution of program

Incremental Parallelization

- Sequential program a special case of a shared-memory parallel program
- Parallel shared-memory programs may only have a single parallel loop
- Incremental parallelization: process of converting a sequential program to a parallel program a little bit at a time

Shared-memory Model vs. Message-passing Model (#2)

- Shared-memory model
 - Execute and profile sequential program
 - Incrementally make it parallel
 - Stop when further effort not warranted
- Message-passing model
 - Sequential-to-parallel transformation requires major effort
 - Transformation done in one giant step rather than many tiny steps

Parallel for Loops

- Fortran programs often express data-parallel operations as `do` loops

```
do i=1,nx, bs
  a(i) = b(i)+c(i)
enddo
```

- OpenMP makes it easy to indicate when the iterations of a loop may execute in parallel
- Compiler takes care of generating code that forks/joins threads and allocates the iterations to threads

Compiler directive

- Compiler directive in Fortran
- A way for the programmer to communicate with the compiler
- Compiler free to ignore directives
- Syntax:

```
!$OMP <rest of directive>
```

```
!$OMP& continuation line
```

```
!$ INTEGER :: i !something only in the  
parallel version of the program
```

Pragma in C or C++

- Syntax:

```
#pragma omp <rest of directive>
```

Parallel do directive

- Fortran:

```
!$OMP parallel do
do i=1,10
    a(i) = b(i) + c(i)
end do
!$OMP end parallel do
```

- Compiler must be able to verify the run-time system will have information it needs to schedule loop iterations

- C

```
#pragma omp for
for (i=0;i<10;i++){
    a[i] = b[i] + c[i];
}
```

Execution Context


- Every thread has its own execution context
- Execution context: address space containing all of the variables a thread may access
- Contents of execution context:
 - static variables
 - dynamically allocated data structures in the heap
 - variables on the run-time stack
 - additional run-time stack for functions invoked by the thread

Shared and Private Variables

- Shared variable: has same address in execution context of every thread
- Private variable: has different address in execution context of every thread
- A thread cannot access the private variables of another thread

Shared and Private Variables

```
int main (int argc, char *argv[])  
{  
    int b[3];  
    char *cptr;  
    int i;  
  
    cptr = malloc(1);  
    #pragma omp parallel for  
    for (i = 0; i < 3; i++)  
        b[i] = i;
```



Using the GCC compilers with OpenMP

- `gfortran -fopenmp -g basic.f90`
- `gcc -fopenmp -g basic.c`
- `setenv OMP_NUM_THREADS 16` if `tcsh/csh`
- `export OMP_NUM_THREADS=16` if `bash/sh`
- `./a.out`

Using the Intel compilers with OpenMP

- `ifort -openmp -g basic.f90`
- `icc -openmp -g basic.c`
- `setenv OMP_NUM_THREADS 16` if `tcsh/csh`
- `export OMP_NUM_THREADS=16` if `bash/sh`
- `./a.out`

Stream example

- Python notebook demo

Function `omp_get_num_procs`

- Returns number of physical processors available for use by the parallel program
- C: **`int omp_get_num_procs(void);`**
- Fortran:

interface

```
function omp_get_num_procs ()  
    use omp_lib_kinds  
    integer ( kind=omp_integer_kind ) :: omp_get_num_procs  
end function omp_get_num_procs
```

end interface

Function `omp_get_thread_num`

- This function returns the thread identification number
- If there are t threads, the ID numbers range from 0 to $t-1$
- The master thread has ID number 0

```
integer omp_get_thread_num ()
```

single Directive

- Suppose we only want to see the output once
- The **single** directive directs compiler that only a single thread should execute the block of code the directive precedes
- Syntax:

```
!$omp single
```

Use of single Pragma

```
!$omp parallel private(i,j)
Do i = 1, m
    low = a(i)
    high = b(i)
    if (low > high) then
        !$OMP single
        write(*,*) "Exiting", I
        !$OMP end single
        exit
    endif
End do
!$omp do
do j = low, high
    c(j) = (c(j) - a(i))/b(i)
End do
!$omp end do
!$omp end parallel
```


Master directive

- Code is executed by the master thread
- Otherwise very similar to the single directive

Function `omp_set_num_threads`

- Uses the parameter value to set the number of threads to be active in parallel sections of code
- May be called at multiple points in a program

```
subroutine omp_set_num_threads (t)  
integer, intent(in) :: t
```

```
void omp_set_num_threads(int num_threads);
```

Hello World - OpenMP

- See ipython notebook

Declaring Private Variables

```
do j = 1,nj
  do i = 1, ni
    a(i,j) = min(a(i,j), a(i,j)+tmp)
  end do
end do
```

- Either loop could be executed in parallel
- We prefer to make outer loop parallel, to reduce number of forks/joins
- We then must give each thread its own private copy of variable **i**

private Clause

- Clause: an optional, additional component to a pragma
- Private clause: directs compiler to make one or more variables private

```
private ( <variable list> )
```

Example Use of private Clause

```
!$OMP parallel do private(i)
DO j = 1, nj
    DO i = 1, ni
        a(i,j) = MIN(a(i,j), a(i,k)+tmp)
    END DO
END DO
```

```
#pragma omp for private(i)
for (j=0; j<nj; j++)
    for (i=0; i<ni; i++) {
        a[j,i] = a[j,i]+tmp
    }
```

firstprivate Clause

- Used to create private variables having initial values identical to the variable controlled by the master thread as the loop is entered
- Variables are initialized once per thread, not once per loop iteration
- If a thread modifies a variable's value in an iteration, subsequent iterations will get the modified value

Use of firstprivate

program **wrong**

```
I = 10
!$OMP PARALLEL PRIVATE(I)
I= I + 1
print *, I
!$OMP END PARALLEL
```

program **correct**

```
I = 10
!$OMP PARALLEL FIRSTPRIVATE(I)
I= I + 1
print *, I
!$OMP END PARALLEL
```


lastprivate Clause

- Sequentially last iteration: iteration that occurs last when the loop is executed sequentially
- **lastprivate** clause: used to copy back to the master thread's copy of a variable the private copy of the variable from the thread that executed the sequentially last iteration

Numerical Integration

$$\pi = \int_0^1 \frac{4.0}{1+x^2} dx$$

$$\sum_{i=0}^N f(x_i) dx_i \approx \pi$$

Critical Sections

```
integer n, i
double precision w, x, area,
pi, f, a
```

```
! function to integrate
f(a) = 4.d0 / (1.d0 + a*a)
```

```
w = 1.0d0/n
area = 0.0d0
```

```
do i = 1, n
    x = w * (i - 0.5d0)
    area = area + f(x)
end do
pi = w * area
```

Race Condition (cont.)

- If we simply parallelize the loop...

```
area = 0.0d0

!$OMP PARALLEL DO PRIVATE(x)
do i = 1, n
    x = w * (i - 0.5d0)
    area = area + f(x)
end do
!$OMP END PARALLEL DO
pi = w * area
```

Race Condition (cont.)

- ... we set up a race condition in which one process may “race ahead” of another and not see its change to shared variable **area**

area

15.230

Answer should be 18.995

Thread A


15.432

Thread B

15.230

area = area + 4.0 / (1.0 + x*x)

Race Condition Time Line



critical directive

- Critical section: a portion of code that only thread at a time may execute
- We denote a critical section by putting the directive

`!$OMP critical`

in front of a block of Fortran code

Correct, But Inefficient, Code

```
! function to integrate
f(a) = 4.d0 / (1.d0 + a*a)

w = 1.0d0/n
area = 0.0d0
!$OMP PARALLEL DO PRIVATE(x) , SHARED(w,n)
do i = 1, n
    x = w * (i - 0.5d0)
    !$OMP CRITICAL
    area = area + f(x)
    !$OMP END CRITICAL
end do
!$OMP END PARALLEL DO
pi = w * area
```


Source of Inefficiency

- Update to **area** inside a critical section
- Only one thread at a time may execute the statement; i.e., it is sequential code
- Time to execute statement significant part of loop
- By Amdahl's Law speedup will be severely constrained

Reductions

- Reductions are so common that OpenMP provides support for them
- May add reduction clause to **parallel do** directive
- Specify reduction operation and reduction variable
- OpenMP takes care of storing partial results in private variables and combining partial results after the loop

reduction Clause

- The reduction clause has this syntax:
reduction (<op> :<variable>)
- Operators
 - +
 - *
 - -
 - .AND.
 - .OR.
 - .EQV.
 - .NEQV.
 - MAX
 - MIN
 - IAND
 - IOR
 - IEOR

π -finding Code with Reduction Clause

```
! function to integrate
f(a) = 4.d0 / (1.d0 + a*a)

w = 1.0d0/n
area = 0.0d0
!$OMP PARALLEL DO PRIVATE(X) , SHARED(w,n)
!$OMP& REDUCTION(+:area)
do i = 1, n
    x = w * (i - 0.5d0)
    area = area + f(x)
end do
!$OMP END PARALLEL DO
pi = w * area
```

Summary

- OpenMP an API for shared-memory parallel programming
- Shared-memory model based on fork/join parallelism
- Data parallelism
 - parallel for pragma
 - reduction clause