


# MPI

## MPI programming

# Flynn's Taxonomy

- Instruction stream
- Data stream
- Single vs. multiple
- Four combinations
  - SISD
  - SIMD
  - MISD
  - MIMD

# Why is our simple distributed-memory machine MIMD?



# SPMD

- SPMD (single program, multiple data): all processors execute same program, but each operates on different portion of problem data
- Easier to program than true MIMD and more flexible than SIMD
- Although most parallel computers today are MIMD architecturally, they are usually programmed in SPMD style

# Message passing

- Most natural and efficient paradigm for distributed-memory systems
- Two-sided, **send** and **receive** communication between processes
- Efficiently portable to shared-memory or almost any other parallel architecture:  
“assembly language of parallel computing” due to universality and detailed, low-level control of parallelism

# More on message passing

- Provides natural synchronization among processes (through blocking receives, for example), so explicit synchronization of memory access is unnecessary
- Sometimes deemed tedious and low-level, but thinking about locality promotes
  - good performance,
  - scalability,
  - portability
- Dominant paradigm for developing portable and scalable applications for massively parallel systems

# Programming a distributed-memory computer

- MPI (Message Passing Interface)
- Message passing standard, universally adopted library of communication routines callable from C, C++, Fortran, (Python)
- 125+ functions—I will introduce a small subset of functions

# MPI-1

- MPI was developed in two major stages, MPI-1 and MPI-2
- Features of MPI-1 include
  - point-to-point communication
  - collective communication process
  - groups and communication domains
  - virtual process topologies
  - environmental management and inquiry
  - profiling interface bindings for Fortran and C



# MPI-2

- Additional features of MPI-2 include:
  - dynamic process management input/output
  - one-sided operations for remote memory access (update or interrogate)
  - memory access bindings for C++
- We will cover no MPI-2

# MPI programs use SPMD model

- Same program runs on each process
- Build executable and link with MPI library
- User determines number of processes and on which processors they will run

# Programming in MPI

use mpi

```
Integer :: ierr  
call MPI_init(ierr)  
.  
.  
.  
call MPI_Finalize(ierr)
```

#include "mpi.h"

```
int ierr;  
ierr = MPI_Init(&argc, &argv);  
.  
.  
.  
ierr = MPI_Finalize();
```

C returns error codes as function values,  
Fortran requires arguments (ierr)

# Programming in MPI

```
use mpi  
integer ierr
```

```
call MPI_init(ierr)  
call MPI_COMM_RANK( MPI_COMM_WORLD, myid, ierr )  
call MPI_COMM_SIZE( MPI_COMM_WORLD, numprocs, ierr )  
.  
.  
.  
call MPI_Finalize(ierr)
```

Determine process id or *rank* (here = myid)  
And number of processes (here = numprocs)

# Determine the processor running on

- `lerr = MPI_Get_processor_name(proc_name, &length);`

# MPI\_COMM\_WORLD

- Is a *communicator*
- Predefined in MPI
- Consists of all processes running at start of program execution
- Process rank and number of processors determined from MPI\_COMM\_WORLD
- Possible to create new communicators

# MPI Hello world


- Write a program similar to the OMP hello world
- Output should be:
  - Hello World from process = %d on node %s
  - Number of mpi processes = %d

# Compiling and Running an MPI program

- See `mpi_intro.ipynb`



# Sending data in a ring



# Example program

- Summing numbers on ring of processors  
initially,  $N$  single numbers per processor
- If I am processor  $myid$ ,
- Store my number in  $x(1:n)$
- For number of steps =  $numprocs - 1$ 
  - Send my  $n$  numbers to process  $myid + 1 \pmod{numprocs}$
  - Receive  $N$   $x$  from process  $myid - 1 \pmod{numprocs}$
  - Once all values have been received, sum  $x(1) + \dots + x(numprocs)$

# Blocking send

- call MPI\_SEND(  
    message,                   e.g., my\_partial\_sum,  
    count,                    number of values in msg  
    data\_type,                e.g., MPI\_DOUBLE\_PRECISION,  
    destination,              e.g., myid + 1  
    tag,                      some info about msg, e.g., store it  
    communicator,            e.g., MPI\_COMM\_WORLD,  
    ierr  
)

All arguments are inputs.

# Fortran MPI Data Types

MPI\_CHARACTER  
MPI\_COMPLEX, MPI\_COMPLEX8, also 16 and 32  
MPI\_DOUBLE\_COMPLEX  
MPI\_DOUBLE\_PRECISION  
MPI\_INTEGER  
MPI\_INTEGER1, MPI\_INTEGER2, also 4 and 8  
MPI\_LOGICAL  
MPI\_LOGICAL1, MPI\_LOGICAL2, also 4 and 8  
MPI\_REAL  
MPI\_REAL4, MPI\_REAL8, MPI\_REAL16

Numbers = numbers of bytes  
Somewhat different in C—see text or Google it

# C MPI Datatypes

MPI_CHAR	8-bit character
MPI_DOUBLE	64-bit floating point
MPI_FLOAT	32-bit floating point
MPI_INT	32-bit integer
MPI_LONG	32-bit integer
MPI_LONG_DOUBLE	64-bit floating point
MPI_LONG_LONG	64-bit integer
MPI_LONG_LONG_INT	64-bit integer
MPI_SHORT	16-bit integer
MPI_SIGNED_CHAR	8-bit signed character
MPI_UNSIGNED	32-bit unsigned integer
MPI_UNSIGNED_CHAR	8-bit unsigned character
MPI_UNSIGNED_LONG	32-bit unsigned integer
MPI_UNSIGNED_LONG_LONG	64-bit unsigned integer
MPI_UNSIGNED_SHORT	16-bit unsigned integer
MPI_WCHAR	Wide (16-bit) unsigned character

# Blocking?

- MPI\_send
  - does not return until the message data and envelope have been buffered in matching receive buffer or temporary system buffer.
  - can complete as soon as the message was buffered, even if no matching receive has been executed by the receiver.
  - MPI buffers or not, depending on availability of space
  - **non-local**: successful completion of the send operation may depend on the occurrence of a matching receive.

# Blocking receive

- call MPI\_RECV(  
    message,                   e.g., my\_partial\_sum,  
    count,                    number of values in msg  
    data\_type,                e.g., MPI\_DOUBLE\_PRECISION,  
    source,                  e.g., myid - 1  
    tag,                    some info about msg, e.g., store it  
    communicator,            e.g., MPI\_COMM\_WORLD,  
    status,                  info on size of message received  
    ierr  
)

# The arguments

- outputs: message, status
- $\text{count} \times \text{size of data\_type}$  determines size of receive buffer:
  - too large message received gives error,
  - too small message is ok
- status must be decoded if needed (MPI\_Get\_Count)



# Blocking receive

- Process must wait until message is received to return from call.
- Stalls progress of program BUT
  - blocking sends and receives enforce process synchronization
  - so enforce consistency of data

# Our program

```
integer ierr (and other dimension statements)
include "mpi.h"
call MPI_init(ierr), MPI_COMM_RANK, MPI_COMM_SIZE
< Processor myid has x(1), x(2) to begin>
count = 1
do j = 1, numprocs-1
    call MPI_send(x(count), 2, ..., mod(myid+1, numprocs), ...)
    count = count + 2
    call MPI_recv(x(count), 2, ..., mod(myid-1, numprocs), ...)
enddo
print*, 'here is my answer', sum(x)
Call MPI_finalize(ierr)
```

# Point-to-Point Communication Modes

## Standard Mode:

### blocking:

MPI\_SEND (buf, count, datatype, dest, tag, comm, ierr)

MPI\_RECV (buf, count, datatype, source, tag, comm,  
status, ierr)

Generally **ONLY** use if you cannot call earlier **AND** there is no other work that can be done!

Standard **ONLY** states that buffers can be used once calls return. It is implementation dependent on when blocking calls return.

Blocking sends **MAY** block until a matching receive is posted. This is not required behavior, but the standard does not prohibit this behavior either. Further, a blocking send may have to wait for system resources such as system managed message buffers.

**Be VERY careful of deadlock when using blocking calls!**

# Requirements for Point to Point Communications

- For a communication to succeed:
  - Sender must specify a valid destination rank.
  - Receiver must specify a valid source rank.
  - The communicator must be the same.
  - Tags must match.
  - Message data types must match.
  - Receiver's buffer must be large enough.

# Wildcarding

- Receiver can wildcard.
- To receive from any source
  - source = MPI\_ANY\_SOURCE
- To receive from any tag
  - tag = MPI\_ANY\_TAG
- Actual source and tag are returned in the receiver's status parameter.

# Communication Envelope

- Envelope information is returned from MPI\_RECV in **status**.
- C:
  - status.MPI\_SOURCE
  - status.MPI\_TAG
  - count via MPI\_Get\_count()
- Fortran:
  - status(MPI\_SOURCE)
  - status(MPI\_TAG)
  - count via MPI\_GET\_COUNT()

# Deadlock

- Deadlock: process waiting for a condition that will never become true
- Easy to write send/receive code that deadlocks
  - Two processes: both receive before send
  - Send tag doesn't match receive tag
  - Process sends message to wrong destination process

# MPI\_ISEND (buf, cnt, dtype, dest, tag, comm, request, ierr)

- Same syntax as MPI\_SEND with the addition of a request handle
- Request is a handle (int in Fortran; MPI\_Request in C) used to check for completeness of the send
- This call returns immediately
- Data in `buf` may not be accessed until the user has completed the send operation
- The send is completed by a successful call to MPI\_TEST or a call to MPI\_WAIT



# MPI\_IRecv(buf, cnt, dtype, source, tag, comm, request, ierr)

- Same syntax as MPI\_RECV except status is replaced with a request handle
- Request is a handle (int in Fortran MPI\_Request in C) used to check for completeness of the recv
- This call returns immediately
- Data in buf may not be accessed until the user has completed the receive operation
- The receive is completed by a successful call to MPI\_TEST or a call to MPI\_WAIT

# MPI\_WAIT (request, status, ierr)

- Request is the handle returned by the non-blocking send or receive call
- Upon return, status holds source, tag, and error code information
- This call does not return until the non-blocking call referenced by *request* has completed
- Upon return, the request handle is freed
- If *request* was returned by a call to MPI\_ISEND, return of this call indicates nothing about the destination process

# MPI\_WAITANY (count, requests, index, status, ierr)

- Requests is an array of handles returned by non-blocking send or receive calls
- Count is the number of requests
- This call does not return until a non-blocking call referenced by one of the *requests* has completed
- Upon return, index holds the index into the array of requests of the call that completed
- Upon return, status holds source, tag, and error code information for the call that completed
- Upon return, the request handle stored in requests[index] is freed

# MPI\_WAITALL (count, requests, statuses, ierr)

- *requests* is an array of handles returned by non-blocking send or receive calls
- *count* is the number of requests
- This call does not return until all non-blocking call referenced by *requests* have completed
- Upon return, *statuses* hold source, tag, and error code information for all the calls that completed
- Upon return, the request handles stored in *requests* are all freed

# MPI\_TEST (request, flag, status, ierr)

- *request* is a handle returned by a non-blocking send or receive call
- Upon return, *flag* will have been set to true if the associated non-blocking call has completed. Otherwise it is set to false
- If *flag* returns true, the request handle is freed and *status* contains source, tag, and error code information
- If *request* was returned by a call to MPI\_ISEND, return with *flag* set to true indicates nothing about the destination process

# MPI\_TESTANY (count, requests, index, flag, status, ierr)

- *requests* is an array of handles returned by non-blocking send or receive calls
- *count* is the number of requests
- Upon return, *flag* will have been set to true if one of the associated non-blocking call has completed. Otherwise it is set to false
- If *flag* returns true, *index* holds the index of the call that completed, the request handle is freed, and *status* contains source, tag, and error code information

# MPI\_TESTALL (count, requests, flag, statuses, ierr)

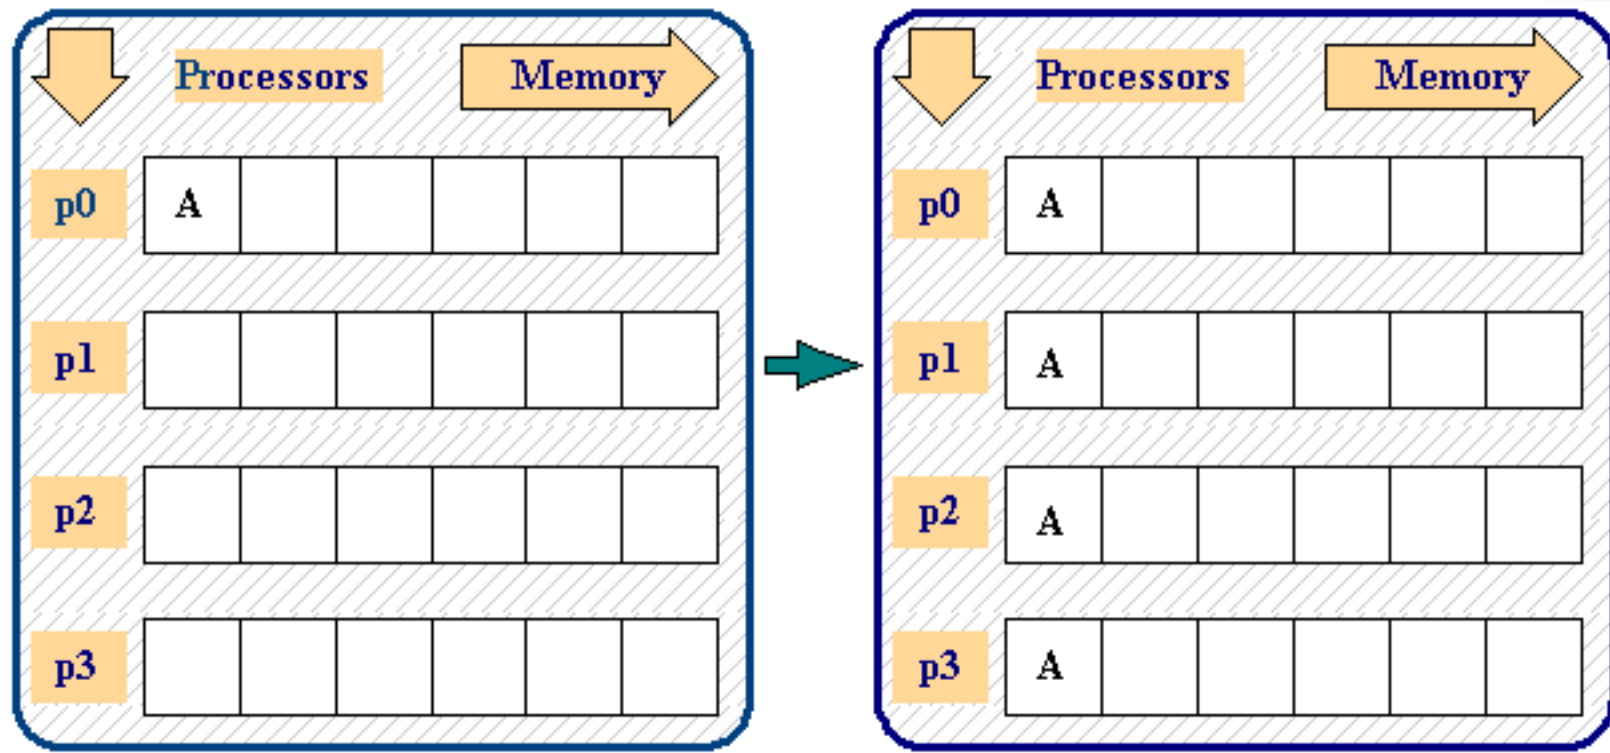
- *requests* is an array of handles returned by non-blocking send or receive calls
- *count* is the number of requests
- Upon return, *flag* will have been set to true if ALL of the associated non-blocking call have completed. Otherwise it is set to false
- If *flag* returns true, all the request handles are freed, and *statuses* contains source, tag, and error code information for each operation

# Collective communication

- One-To-All
  - MPI\_Bcast(), MPI\_Scatter(), MPI\_Scatterv()
- All-To-One
  - MPI\_Gather(), MPI\_Gatherv(), MPI\_Reduce()
- All-To-All
  - MPI\_Allgather(), MPI\_Allgatherv(), MPI\_Allreduce()
- Other
  - MPI\_Barrier()




# Broadcast



```
send_count = 1;  
root = 0;  
MPI_Bcast ( &a, send_count, MPI_INT, root, comm )
```

Figure from MPI-tutor: <http://www.citutor.org/index.php>

# Reduction



```
count = 1;
```

```
rank = 0;
```


```
MPI_Reduce ( &a, &x, count, MPI_REAL, MPI_SUM, rank, MPI_COMM_WORLD );
```

Figure from MPI-tutor: <http://www.citutor.org/index.php>

# Reduction operations

Operation	Description
<b>MPI_MAX</b>	maximum
<b>MPI_MIN</b>	minimum
<b>MPI_SUM</b>	sum
<b>MPI_PROD</b>	product
<b>MPI LAND</b>	logical and
<b>MPI_BAND</b>	bit-wise and
<b>MPI_LOR</b>	logical or
<b>MPI BOR</b>	bit-wise or
<b>MPI_LXOR</b>	logical xor
<b>MPI_BXOR</b>	bitwise xor
<b>MPI_MINLOC</b>	computes a global minimum and an index attached to the minimum value -- can be used to determine the rank of the process containing the minimum value
<b>MPI_MAXLOC</b>	computes a global maximum and an index attached to the rank of the process containing the minimum value

# Gather



```
send_count = 1;  
recv_count = 1;  
recv_rank = 0;  
MPI_Gather ( &a, send_count, MPI_REAL, &a, recv_count, MPI_REAL, recv_rank,  
MPI_COMM_WORLD );
```

# All-gather

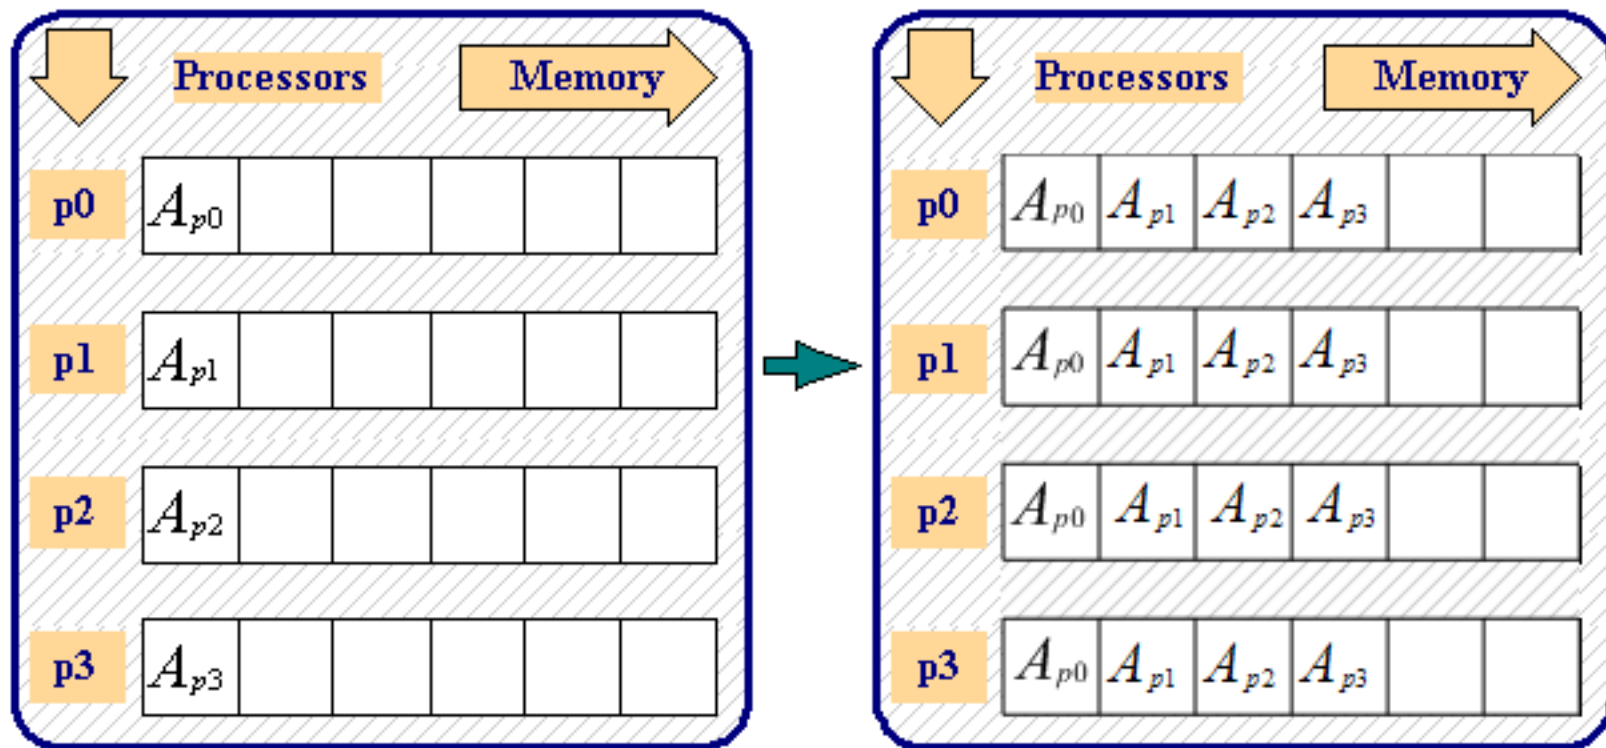
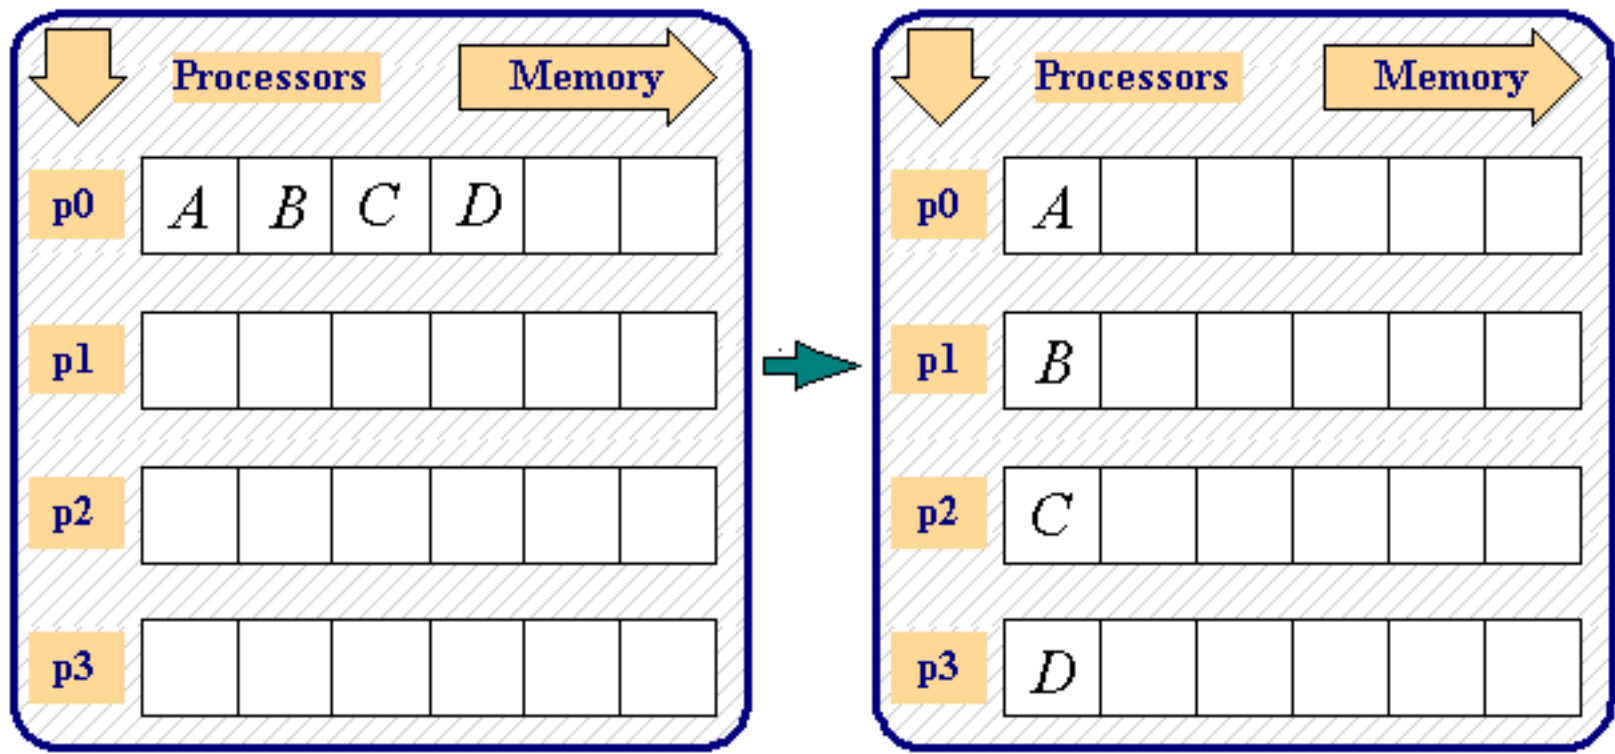


Figure from MPI-tutor: <http://www.citutor.org/index.php>

# Scatter



```
recv_count = 1;  
send_rank = 0;  
MPI_Scatter ( &a, send_count, MPI_REAL,  
             &a, recv_count, MPI_REAL,  
             send_rank, MPI_COMM_WORLD );
```

# Question 1

- You want to do a simple broadcast of variable `abc[7]` in processor 0 to the same location in all other processors of the communicator. What is the correct syntax of the call to do this broadcast?
  1. `MPI_Bcast ( &abc[7], 1, MPI_REAL, 0, comm )`
  2. `MPI_Bcast ( &abc, 7, MPI_REAL, 0, comm )`
  3. `MPI_Broadcast ( &abc[7], 1, MPI_REAL, 0, comm )`

# Summary

- MPI is the standard for distributed parallel programming
- Best approach is probably hybrid
  - MPI for inter node communication
  - OpenMP for and other directives for parallelism within a node
- If possible use existing libraries
  - Global Arrays <http://hpc.pnl.gov/globalarrays/>
  - PETSc <http://www.mcs.anl.gov/petsc/>